
A Latent Neural ODE-VAE for Modeling Hippocampal Population Activity on Low-Dimensional Manifolds

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Neural activity can be represented as a trajectory in a high-dimensional space where
2 each axis corresponds to one neuron. A central question in systems neuroscience
3 is whether these trajectories are constrained to low-dimensional manifolds that
4 jointly encode external variables and trial-to-trial internal variability. Building on
5 this question, we develop and analyze a latent Neural ODE variational autoencoder
6 (ODE-VAE) for calcium population recordings. The model encodes trial onset ac-
7 tivity into a stochastic latent initial condition, evolves it through a continuous-time
8 mixture-of-experts neural ODE, and decodes latent trajectories back to neural activi-
9 ty. Relative to earlier ODE-VAE versions, v5 adds transition-aware reconstruction
10 regularization and a soft locally linear embedding (LLE) latent constraint. Using
11 repository-grounded experiments, we report strong reconstruction on synthetic
12 random-foraging data (best seed $R^2 = 0.9789$) and mixed outcomes on the E65
13 hippocampal dataset, including an archived ODE-VAE run at $R^2 = 0.4368$, strict
14 raw-space v5 sweeps with best-epoch $R^2 \leq 0.0569$, and a no-PCA v6 variant
15 at $R^2 = 0.0956$. These results show that continuous-time latent manifold mod-
16 els can capture structured dynamics, but evaluation protocol, preprocessing, and
17 regularization balance strongly affect performance on real neural data.

18

1 Introduction

19 Neural activity can be described as a point in a high-dimensional coordinate system, where each
20 coordinate axis represents a single neuron’s activity. Prior work shows that population trajectories
21 often occupy low-dimensional manifolds, in motor cortex, hippocampus, and entorhinal circuits
22 [Gallego et al., 2017, Russo et al., 2018, O’Keefe and Dostrovsky, 1971, Aronov et al., 2017, Nieh
23 et al., 2021]. This geometric view supports the broader “cognitive map” hypothesis: neural circuits
24 organize both physical and abstract knowledge into structured latent spaces [Tolman, 1948, O’Keefe
25 and Nadel, 1978, Stachenfeld et al., 2017, Bellmund et al., 2018].

26 Recent manifold analyses of hippocampal recordings report intrinsic dimensionalities around 4–6 and
27 show that manifold geometry can explain both task variables and trial-level variability [Low et al.,
28 Chaudhuri et al., 2019, Nieh et al., 2021]. However, many pipelines are multi-stage: fit local
29 transition models, build pairwise distances, embed points, and separately map back to activity [Low
30 et al., 2018, Tenenbaum et al., 2000, Yu et al., 2009]. This decomposition gives useful geometry
31 but does not explicitly enforce globally smooth continuous-time dynamics in a single end-to-end
32 objective.

33 This paper develops a continuous-time generative alternative tailored to the repository’s ODE-VAE
34 codebase. Starting from an initial latent ODE-VAE, we focus on v5, which adds (i) a mixture-of-
35 experts (MoE) latent vector field, (ii) transition-aware reconstruction regularization, and (iii) soft LLE

36 regularization in latent space. We evaluate this family on synthetic and real hippocampal calcium
37 data and document where the method works and where it fails.

38 **Contributions.**

- 39 • We present a mathematically explicit latent Neural ODE-VAE formulation for trial-based
40 neural population sequences, including MoE latent dynamics and local-geometric regulariz-
41 ers.
42 • We instantiate this formulation using the repository’s v1–v6 implementations, with v5 as
43 the main model and v6 as a no-PCA extension.
44 • We provide artifact-grounded empirical analysis on synthetic and E65 datasets, including
45 seed sensitivity and configuration-dependent instability.
46 • We identify concrete protocol factors (split strategy, landmark subsampling, and metric
47 space) that materially influence reported R^2 .

48 **2 Related Work**

49 Our approach sits at the intersection of hippocampal manifold neuroscience and deep latent dynam-
50 ical systems. In neuroscience, place cells and cognitive map theories motivate low-dimensional
51 geometric structure in population activity [O’Keefe and Dostrovsky, 1971, O’Keefe and Nadel, 1978,
52 Eichenbaum and Cohen, 2014]. More recent work extends these ideas to abstract and non-spatial
53 variables [Constantinescu et al., 2016, Schuck and Niv, 2019, Park et al., 2020, Aronov et al., 2017,
54 Nieh et al., 2021]. Population-level manifold analyses in hippocampus and other circuits show that
55 latent geometry can explain behavior and internal state variability [Low et al., 2018, Chaudhuri et al.,
56 2019, Gallego et al., 2017].

57 In machine learning, variational autoencoders [Kingma and Welling, 2014] and neural ODEs [Chen
58 et al., 2018] enable continuous-time latent-variable models, including latent ODEs for irregular
59 trajectories [Rubanova et al., 2019]. Our model adopts this framework but targets neuroscientific
60 interpretability: trialized sequences, manifold visualizations, and explicit regularizers for smoothness,
61 transition consistency, and local linearity.

62 **3 Problem Setup and Data**

63 Let $x_b(t_\ell) \in \mathbb{R}^N$ denote population activity for trial $b \in \{1, \dots, B\}$ at resampled time index
64 $\ell \in \{1, \dots, L\}$.

65 **E65 dataset.** The repository’s main real dataset is E65_data.npz, containing calcium activity
66 ($N = 375$, $T = 7434$), trial IDs, timestamps, and aligned behavioral covariates. In the v5 prepro-
67 cessing path: (i) PCA retains 95% variance, producing $K = 129$ components; (ii) trials are resampled
68 to $L = 120$ bins using trial_len_s=12 and fps=10; (iii) first 10 trials are dropped, yielding 180
69 usable trials; (iv) default validation holds out 3 trials (train 177 / val 3).

70 **Synthetic benchmark.** The repository also includes synthetic_rat_data.npz (4000 frames,
71 300 neurons, 20 trials), used to test recoverability of smooth low-dimensional dynamics under
72 controlled conditions.

73 **4 Model: Latent Neural ODE-VAE**

74 **4.1 Stochastic encoder**

75 For each trial, the encoder uses only $x_b(t_1)$ and outputs a diagonal Gaussian posterior on the latent
76 initial state:

$$q_\phi(z_{0,b} | x_b(t_1)) = \mathcal{N}(\mu_b, \text{diag}(\sigma_b^2)), \quad (1)$$

77 with reparameterization

$$z_{0,b} = \mu_b + \sigma_b \odot \epsilon, \quad \epsilon \sim \mathcal{N}(0, I). \quad (2)$$

78 **4.2 Continuous-time latent dynamics**

79 Latent trajectories are generated by a neural ODE:

$$\frac{dz_b(t)}{dt} = f_\theta(z_b(t), t), \quad z_b(t_1) = z_{0,b}. \quad (3)$$

80 In v5, f_θ is a mixture of experts:

$$f_\theta(z, t) = \sum_{e=1}^E \pi_e(z) f_e(z), \quad \pi(z) = \text{softmax}(g(z)), \quad (4)$$

81 with $E = 4$ latent experts by default and Dormand–Prince integration (`dopri5`).

82 **4.3 Decoder family**

83 A decoder maps latent states back to observations:

$$\hat{x}_b(t_\ell) = g_\psi(z_b(t_\ell)). \quad (5)$$

84 The codebase supports MLP, neuron-aware, local-attention, and MoE decoders; v5 default is MoE
85 decoder with 8 decoder experts.

86 **5 Training Objective and Regularization**

87 The base objective combines reconstruction and KL terms:

$$\mathcal{L}_{\text{base}} = \mathcal{L}_{\text{rec}} + \beta \mathcal{L}_{\text{KL}}, \quad (6)$$

88 where

$$\mathcal{L}_{\text{rec}} = \frac{1}{BLK} \sum_{b,\ell} \|\hat{x}_b(t_\ell) - x_b(t_\ell)\|_2^2, \quad (7)$$

$$\mathcal{L}_{\text{KL}} = \frac{1}{B} \sum_b D_{\text{KL}}(q_\phi(z_{0,b} | x_b(t_1)) \| \mathcal{N}(0, I)). \quad (8)$$

Smoothness regularization.

$$\mathcal{L}_{\text{smooth}} = \frac{1}{B(L-1)D} \sum_{b,\ell} \left\| \frac{z_b(t_{\ell+1}) - z_b(t_\ell)}{t_{\ell+1} - t_\ell} \right\|_2^2. \quad (9)$$

Transition-aware regularization (v5).

$$\mathcal{L}_{\text{trans}} = \frac{1}{B(L-1)K} \sum_{b,\ell} \left\| (\hat{x}_b(t_{\ell+1}) - \hat{x}_b(t_\ell)) - (x_b(t_{\ell+1}) - x_b(t_\ell)) \right\|_2^2. \quad (10)$$

90 This term is linearly warmed up for the first 30 epochs.

91 **Soft LLE latent regularization (v5).** For flattened latent points $\{z_i\}_{i=1}^M$, with k -NN set $\mathcal{N}_k(i)$:

$$\mathcal{L}_{\text{LLE}} = \frac{1}{M} \sum_{i=1}^M \left\| z_i - \sum_{j \in \mathcal{N}_k(i)} w_{ij} z_j \right\|_2^2, \quad w_{ij} \propto \exp\left(-\frac{\|z_i - z_j\|_2}{\tau}\right). \quad (11)$$

92 Default parameters: $k = 8$, $M \leq 256$, $\tau = 0.1$.

Total loss.

$$\mathcal{L} = \mathcal{L}_{\text{rec}} + \beta_t \mathcal{L}_{\text{KL}} + \lambda_{\text{smooth}} \mathcal{L}_{\text{smooth}} + \lambda_{\text{trans},t} \mathcal{L}_{\text{trans}} + \lambda_{\text{LLE}} \mathcal{L}_{\text{LLE}}. \quad (12)$$

93 The KL coefficient β_t is warmed up over 30 epochs to a final value $\beta = 0.02$.

Table 1: Synthetic random-foraging benchmark from `seed_sweep_results.txt`.

Seed	Final R^2	Best validation loss
1	0.9789	0.06642
42	0.6757	0.40521
1337	0.9116	0.16298
2025	0.7467	0.31145
777	0.0000 (NaN collapse)	∞
Mean (all seeds)	0.6626	—
Mean (non-collapsed seeds)	0.8282	—

Table 2: E65 run artifacts from saved metadata/checkpoints.

Artifact (path)	Variant summary	Best reported R^2	Final repo
<code>runs/ode_vae_E65/run_metadata.json</code>	archived ODE-VAE	—	0.4368
<code>runs/ode_vae_E65/ld5_b0.02_ls0.0005/run_metadata.json</code>	v5, $d = 5$	0.0354	-0.23
<code>runs/ode_vae_E65/ld5_b0.02_ls0.0002/run_metadata.json</code>	v5, $d = 5$	0.0353	-0.13
<code>runs/ode_vae_E65/ld8_b0.02_ls0.0005/run_metadata.json</code>	v5, $d = 8$	0.0569	-0.20
<code>src/pt_files/final_metrics.pt</code>	v6, no PCA	—	0.0956

94 6 Experimental Protocol

95 6.1 Configurations

96 Main v5 settings from `config.txt`: latent dimension 5 (with sweep to 8), batch size 8, 150
97 epochs, Adam optimizer (learning rate 0.002, weight decay 10^{-5}), $\lambda_{\text{smooth}} \in \{5 \times 10^{-4}, 2 \times 10^{-4}\}$,
98 $\lambda_{\text{trans}} = 0.01$, $\lambda_{\text{LLE}} = 0.01$, landmark count 100, and baseline correction enabled.

99 6.2 Metrics

100 The primary training metric is coefficient of determination,

$$R^2 = 1 - \frac{\sum(X - \hat{X})^2}{\sum(X - \bar{X})^2}. \quad (13)$$

101 The repository contains two evaluation styles: (i) PCA-space R^2 and (ii) strict raw-neuron-space R^2
102 via inverse PCA and de-normalization. We report artifact values as stored in each run metadata file.

103 7 Results

104 7.1 Synthetic benchmark: high ceiling with seed sensitivity

105 Table 1 summarizes the repository seed sweep on synthetic data. Best seed reaches $R^2 = 0.9789$, but
106 one seed collapses to NaN/0.0, indicating optimizer sensitivity.

107 7.2 E65 hippocampal data: mixed performance across variants

108 Table 2 reports all E65 artifact metrics found in the repository for the ODE-VAE family. An archived
109 ODE-VAE run reports $R^2 = 0.4368$. Strict raw-space v5 sweep artifacts are substantially lower
110 (best-epoch $R^2 \leq 0.0569$, negative final R^2 in three runs), while v6 (no PCA) records $R^2 = 0.0956$.

111 7.3 Manifold interpretability

112 The codebase saves latent manifold projections (MDS) and reconstruction diagnostics for each run.
113 Figure 1 shows an example latent trajectory embedding from the trained model artifacts.

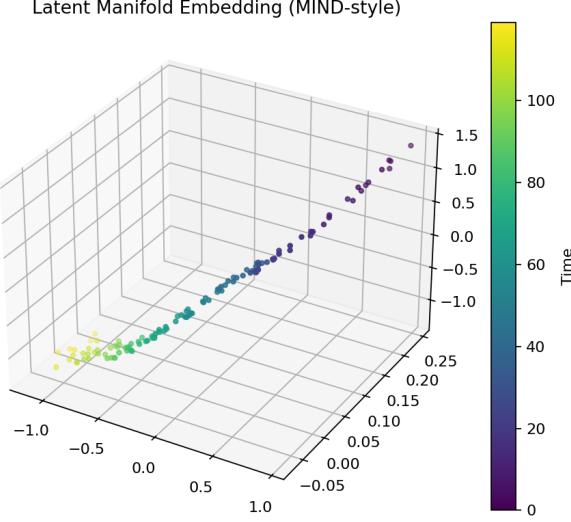


Figure 1: Latent manifold embedding produced by the ODE-VAE analysis pipeline.

114 8 Discussion

115 The model captures the intended inductive bias: low-dimensional continuous latent trajectories
 116 with explicit geometric regularization. On synthetic data, this bias is highly effective. On real E65
 117 recordings, however, results are sensitive to implementation and evaluation choices.

118 Three factors emerge from the repository artifacts:

- 119 1. **Metric-space mismatch.** PCA-space training can look favorable while strict raw-space R^2
 120 may degrade.
- 121 2. **Data-efficiency tradeoff.** Landmark subsampling (100 selected sequences from 180 usable
 122 trials) accelerates training but may reduce generalization.
- 123 3. **Optimization stability.** Strong regularization with small validation sets (3 trials) and stiff
 124 latent dynamics can produce unstable or negative final R^2 , despite early high points.

125 These observations suggest that future gains likely require protocol-level changes in addition to
 126 architectural changes: larger and randomized holdout splits, early stopping on a stable cross-validated
 127 objective, trial-level (not frame-level) landmark selection, and direct raw-space reconstruction losses.

128 9 Limitations and Reproducibility

129 This study is bounded by available run artifacts in the repository and inherits version-specific logging
 130 differences. In particular, some run files report “best” and “final” R^2 under different conditions, and
 131 not all checkpoints include identical metadata fields. We therefore report values exactly as saved in
 132 each artifact path. The codebase also indicates unresolved training fragility (including occasional
 133 NaN collapse), which should be addressed before definitive biological claims.

134 10 Conclusion

135 We presented a mathematically grounded latent Neural ODE-VAE framework for neural manifold
 136 modeling and analyzed the repository’s v1–v6 trajectory with v5 as the primary model. The method
 137 can recover smooth low-dimensional dynamics and high synthetic reconstruction quality, but real-
 138 data performance remains sensitive to preprocessing and evaluation protocol. This work provides
 139 a formalized foundation and a clear set of engineering directions for turning ODE-VAE manifold
 140 modeling into a robust neuroscience analysis tool.

141 **Societal impact.** This work is basic research on neural representation learning from animal neuro-
142 science data and has no immediate direct societal deployment. Potential long-term impact is improved
143 scientific understanding of memory and cognition.

144 **References**

- 145 Dmitriy Aronov, Rachel Nevers, and David W Tank. Mapping of a non-spatial dimension by the
146 hippocampal-entorhinal circuit. *Nature*, 543:719–722, 2017.
- 147 Jacob L S Bellmund, Peter Gardenfors, Edvard I Moser, and Christian F Doeller. Navigating cognition:
148 Spatial codes for human thinking. *Science*, 362:eaat6766, 2018.
- 149 Rishidev Chaudhuri, Burak Gercek, Bikash Pandey, Adrien Peyrache, and Ila Fiete. The intrinsic
150 attractor manifold and population dynamics of a canonical cognitive circuit across waking and
151 sleep. *Nature Neuroscience*, 22:1512–1520, 2019.
- 152 Ricky TQ Chen, Yulia Rubanova, Jesse Bettencourt, and David Duvenaud. Neural ordinary differential
153 equations. *NeurIPS*, 2018.
- 154 Alexandra O Constantinescu, Jill X O'Reilly, and Timothy EJ Behrens. Organizing conceptual
155 knowledge in humans with a grid-like code. *Science*, 352:1464–1468, 2016.
- 156 Howard Eichenbaum and Neal J Cohen. Can we reconcile the declarative memory and spatial
157 navigation views on hippocampal function? *Neuron*, 83:764–770, 2014.
- 158 Juan A Gallego, Matthew G Perich, Lee E Miller, and Sara A Solla. Neural manifolds for the control
159 of movement. *Neuron*, 94:978–984, 2017.
- 160 Diederik P Kingma and Max Welling. Auto-encoding variational bayes. *ICLR*, 2014.
- 161 Ryan J Low, Sean Lewallen, Dmitriy Aronov, Rachel Nevers, and David W Tank. Probing variability
162 in a cognitive map using manifold inference from neural dynamics. *bioRxiv*, 2018.
- 163 Christopher J MacDonald, Kyle Q Lepage, Uri T Eden, and Howard Eichenbaum. Hippocampal time
164 cells bridge the gap in memory for discontiguous events. *Neuron*, 71:737–749, 2011.
- 165 Edward H Nieh et al. Geometry of abstract learned knowledge in the hippocampus. *Nature*, 595:80–
166 84, 2021.
- 167 John O'Keefe and Jonathan Dostrovsky. The hippocampus as a spatial map. *Brain Research*,
168 34:171–175, 1971.
- 169 John O'Keefe and Lynn Nadel. *The Hippocampus as a Cognitive Map*. Clarendon Press, 1978.
- 170 Sang Ah Park, David S Miller, Hamed Nili, Charan Ranganath, and Erie D Boorman. Map making:
171 Constructing, combining, and inferring on abstract cognitive maps. *Neuron*, 107:1226–1238.e8,
172 2020.
- 173 Eftychios A Pnevmatikakis et al. Simultaneous denoising, deconvolution, and demixing of calcium
174 imaging data. *Neuron*, 89:285–299, 2016.
- 175 Eftychios A Pnevmatikakis and Andrea Giovannucci. NoRMCorre: An online algorithm for piecewise
176 rigid motion correction of calcium imaging data. *Journal of Neuroscience Methods*, 291:83–94,
177 2017.
- 178 Stefano Recanatesi et al. Predictive learning as a network mechanism for extracting low-dimensional
179 latent space representations. *Nature Communications*, 12:1417, 2021.
- 180 Yulia Rubanova, Ricky TQ Chen, and David Duvenaud. Latent ordinary differential equations for
181 irregularly-sampled time series. *NeurIPS*, 2019.
- 182 Abigail A Russo et al. Motor cortex embeds muscle-like commands in an untangled population
183 response. *Neuron*, 97:953–966.e8, 2018.

- 184 Nicolas W Schuck and Yael Niv. Sequential replay of nonspatial task states in the human hippocampus.
185 *Science*, 364:eaaw5181, 2019.
- 186 Kimberly L Stachenfeld, Matthew M Botvinick, and Samuel J Gershman. The hippocampus as a
187 predictive map. *Nature Neuroscience*, 20:1643–1653, 2017.
- 188 Joshua B Tenenbaum, Vin de Silva, and John C Langford. A global geometric framework for
189 nonlinear dimensionality reduction. *Science*, 290:2319–2323, 2000.
- 190 Edward C Tolman. Cognitive maps in rats and men. *Psychological Review*, 55:189–208, 1948.
- 191 Byron M Yu et al. Gaussian-process factor analysis for low-dimensional single-trial analysis of neural
192 population activity. *Journal of Neurophysiology*, 102:614–635, 2009.